

Sahil Ashok Vangari

AI Engineer | LLM Systems | Backend (Python)

London, United Kingdom | +447818962058 | sahil.vangari0411@gmail.com | [linkedin.com/in/sahil-vangari](https://www.linkedin.com/in/sahil-vangari) | <https://github.com/SilentXD11>

PROFESSIONAL SUMMARY

Results-driven AI Engineer specialising in LLM-powered systems, scalable backend infrastructure, and production AI deployment. Proven expertise building end-to-end applications integrating computer vision, NLP, and generative models. Excel at latency optimisation, high-performance REST APIs, asynchronous processing, and AWS cloud deployment. Delivered measurable performance improvements and production-ready AI solutions.

TECHNICAL SKILLS

Languages: Python (Advanced), JavaScript, SQL

AI & Machine Learning: Transformers, LLMs, Hugging Face, PyTorch, BERT, NLP, Computer Vision, YOLO, Diffusion Models

Backend: FastAPI, Flask, REST APIs, Async Processing

Cloud & DevOps: AWS (EC2, Lambda, S3, API Gateway), Docker, CI/CD

Databases: PostgreSQL, MongoDB

LEADERSHIP & PRODUCT EXPERIENCE

SIES GST Student Council (Design & Product) | Apr 2022 – Jun 2024

- Led cross-functional teams on end-to-end technical AI projects, improving delivery timelines by 20% through process optimisation and agile workflows
- Mentored team members in product design, technical implementation, and AI best practices, accelerating team capability and output
- Drove projects from concept to production deployment, ensuring alignment between technical delivery and business requirements

KEY PROJECTS

Transformer-Based Fake Review Detection API

BERT, NLP, Flask | Jan - Feb 2026

- Fine-tuned BERT transformer model for real-time fake review and spam detection with high accuracy
- Developed low-latency Flask API capable of processing user-generated content at scale
- Built complete NLP pipeline (preprocessing, tokenisation, class balancing) to handle imbalanced datasets effectively

AI Image Analysis & Generation System

LLM, Computer Vision, AWS, Docker | Jan 2025 - Sept 2025

- Built multi-modal production system combining image captioning, YOLO object detection, and text-to-image generation (diffusion models) via MiniGPT-style pipeline
- Designed asynchronous inference pipeline that dramatically reduced response latency
- Deployed containerised microservices on AWS with full CI/CD pipeline for scalable, reliable production use
- Created unified FastAPI serving multiple models simultaneously

FurSaver – AI Platform (Production MVP)

Flask, PostgreSQL | Jul 2023 - Dec 2023

- Developed backend for real-world AI adoption platform that reached production stage
- Optimised database queries and API structure, reducing response time by 30%
- Implemented secure authentication and scalable REST API architecture supporting growth

Agentic AI Shopping Assistant

LLM Agents, AWS Lambda | Feb 2022 - Jul 2022

- Created autonomous LLM agent using multi-step reasoning and external API integration for personalised recommendations
- Designed fully serverless architecture with AWS Lambda + API Gateway for cost-efficient scaling
- Integrated real-time external APIs to deliver dynamic shopping recommendations

EDUCATION

MSc Artificial Intelligence (Merit)

Queen Mary University of London | Sept 2024 – Sept 2025

CERTIFICATIONS

- Java, Android 12 & Kotlin Development Masterclass – Udemy | AZ900 Microsoft Azure - Udemy
- Career Essentials in Generative AI – Microsoft & LinkedIn
- AI Agent Course – Hugging Face | LangChain & LangSmith – LangChain Academy
- ADK Crash Course – Google | Artificial Intelligence – Great Learning