

# Sahil Ashok Vangari

AI Engineer (LLMs | Backend Systems | Production AI)

London, United Kingdom | +447818962058 | [sahil.vangari0411@gmail.com](mailto:sahil.vangari0411@gmail.com) | [linkedin.com/in/sahil-vangari](https://www.linkedin.com/in/sahil-vangari) | <https://github.com/SilentXD11>

## PROFILE

AI Engineer focused on building **production-grade LLM and multi-modal systems** with strong backend foundations. Experienced in designing **low-latency APIs, asynchronous pipelines, and scalable AI deployments on AWS**. Proven ability to integrate **LLMs, computer vision, and diffusion models** into unified, real-world systems.

---

## TECHNICAL SKILLS

**Languages:** Python (Advanced), JavaScript, SQL

**AI/ML:** Transformers, LLMs, BERT, Hugging Face, PyTorch, YOLOv8, Diffusion Models

**Backend:** FastAPI, Flask, REST APIs, Async Processing

**Cloud/DevOps:** AWS (EC2, Lambda, S3, API Gateway), Docker, CI/CD

**Databases:** PostgreSQL, MongoDB

---

## PROJECT EXPERIENCE

### AI Image Analysis & Generation System (Production-Scale)

*LLMs, Computer Vision, AWS, Docker* | Jan 2025 – Sept 2025

- Built a **multi-modal AI system** combining image captioning, YOLO object detection, and text-to-image generation
- Designed **asynchronous inference pipeline**, reducing response latency by ~40% and improving system throughput
- Developed unified **FastAPI backend serving multiple models concurrently**
- Deployed containerised microservices on AWS with CI/CD pipeline ensuring scalable and reliable performance
- Enabled real-time interaction across multiple AI models via a single API

---

### Transformer-Based Fake Review Detection API

*BERT, NLP, Flask* | Jan 2026 – Feb 2026

- Fine-tuned BERT model achieving **~90% accuracy** on imbalanced dataset
- Built **low-latency API (<200ms response time)** for real-time spam detection
- Implemented complete NLP pipeline (preprocessing, tokenisation, class balancing)
- Designed system for scalable handling of user-generated content
- **GitHub:** [https://github.com/SilentXD11/Fake\\_Review](https://github.com/SilentXD11/Fake_Review)

---

### FurSaver – AI Platform (Production MVP)

*Flask, PostgreSQL* | Jul 2023 – Dec 2023

- Developed backend for production MVP with real-world usage
- Reduced API response time by **30% through query optimisation and efficient schema design**
- Built secure authentication and scalable REST API architecture
- Improved backend performance under increasing user load
- **GitHub:** [https://github.com/SilentXD11/Fursaver\\_main](https://github.com/SilentXD11/Fursaver_main)

---

### Agentic AI Shopping Assistant

*LLM Agents, AWS Lambda* | Feb 2022 – Jul 2022

- Built autonomous LLM agent using **multi-step reasoning and external API integration**
- Designed **serverless architecture (AWS Lambda + API Gateway)** for cost-efficient scaling
- Delivered dynamic, real-time personalised recommendations
- Implemented modular architecture for easy extension of tools and APIs

---

## EDUCATION

### MSc Artificial Intelligence (Merit)

Queen Mary University of London | 2024 – 2025

### Bachelors in Computer Engineering - 8.56

SIES Graduate School of Technology | 2020-2024

---

## CERTIFICATIONS

- Generative AI – Microsoft & LinkedIn | Google ADK Crash Course
- Hugging Face AI Agents Course | Azure AZ-900 (Fundamentals)
- LangChain & LangSmith – LangChain Academy